google / gemma-4-E4B-it

like 926 | Following Google 54.3k

Any-to-Any | Transformers | Safetensors | gemma4 | image-text-to-text | Eval Results | License: apache-2.0

Deploy | Use this model

main | gemma-4-E4B-it / model.safetensors

MaartenGr Preparing for release! 292a7e2

download | Copy download link | history | blame | contribute | delete

Safe

16 GB

This file is stored with Xet . It is too big to display, but you can still download it.

Xet Pointer Details (Raw pointer file)

Xet hash: be9eb4387fbda28c3f67dc5f0e68a3bf381d15e5aa3e1a51559db2f08f931626
Size of remote file: 16 GB · **SHA256:**
cfbd3d2f1cd71bd471c37fe2bf8546d5028d41e5736f64e1ca6c6b8893125503

Xet efficiently stores Large Files inside Git, intelligently splitting files into unique chunks and accelerating uploads and downloads. More info.

Metadata	Value
format	pt

Tensors	Shape
model.embed_audio.embedding_projection.weight	[2 560, 1 536
model.embed_vision.embedding_projection.weight	[2 560, 768]
model.audio_tower (3)	
model.audio_tower.layers (12)	
model.audio_tower.layers.0 (6)	
model.audio_tower.layers.0.feed_forward (2)	
model.audio_tower.layers.0.feed_forward1 (12)	

Tensors	Shape
model.audio_tower.layers.0.feed_forward2 (12)	
model.audio_tower.layers.0.lconv1 (13)	
model.audio_tower.layers.0.norm_out.weight	[1 024]
model.audio_tower.layers.0.norm_post_attn.weight	[1 024]
model.audio_tower.layers.0.norm_pre_attn.weight	[1 024]
model.audio_tower.layers.0.self_attn (6)	
model.audio_tower.layers.0.self_attn.k_proj (5)	
model.audio_tower.layers.0.self_attn.k_proj.input_max	scalar
model.audio_tower.layers.0.self_attn.k_proj.input_min	scalar
model.audio_tower.layers.0.self_attn.k_proj.linear.weight	[1 024, 1 024
model.audio_tower.layers.0.self_attn.k_proj.output_max	scalar
model.audio_tower.layers.0.self_attn.k_proj.output_min	scalar
model.audio_tower.layers.0.self_attn.per_dim_scale	[128]
model.audio_tower.layers.0.self_attn.post (5)	
model.audio_tower.layers.0.self_attn.post.input_max	scalar
model.audio_tower.layers.0.self_attn.post.input_min	scalar
model.audio_tower.layers.0.self_attn.post.linear.weight	[1 024, 1 024
model.audio_tower.layers.0.self_attn.post.output_max	scalar
model.audio_tower.layers.0.self_attn.post.output_min	scalar
model.audio_tower.layers.0.self_attn.q_proj (5)	
model.audio_tower.layers.0.self_attn.q_proj.input_max	scalar
model.audio_tower.layers.0.self_attn.q_proj.input_min	scalar
model.audio_tower.layers.0.self_attn.q_proj.linear.weight	[1 024, 1 024
model.audio_tower.layers.0.self_attn.q_proj.output_max	scalar
model.audio_tower.layers.0.self_attn.q_proj.output_min	scalar
model.audio_tower.layers.0.self_attn.relative_k_proj.weight	[1 024, 1 024
model.audio_tower.layers.0.self_attn.v_proj (5)	
model.audio_tower.layers.0.self_attn.v_proj.input_max	scalar
model.audio_tower.layers.0.self_attn.v_proj.input_min	scalar
model.audio_tower.layers.0.self_attn.v_proj.linear.weight	[1 024, 1 024
model.audio_tower.layers.0.self_attn.v_proj.output_max	scalar
model.audio_tower.layers.0.self_attn.v_proj.output_min	scalar
model.audio_tower.layers.1 (6)	
model.audio_tower.layers.2 (6)	
model.audio_tower.layers.3 (6)	
model.audio_tower.layers.4 (6)	
model.audio_tower.layers.5 (6)	
model.audio_tower.layers.6 (6)	

Tensors	Shape
model.audio_tower.layers.7 (6)	
model.audio_tower.layers.8 (6)	
model.audio_tower.layers.9 (6)	
model.audio_tower.layers.10 (6)	
model.audio_tower.layers.11 (6)	
model.audio_tower.subsample_conv_projection (2)	
model.audio_tower.subsample_conv_projection.input_proj_linear.weight	[1 024, 1 024
model.audio_tower.subsample_conv_projection.layer (2)	
model.audio_tower.subsample_conv_projection.layer0 (2)	
model.audio_tower.subsample_conv_projection.layer1 (2)	
model.audio_tower.output_proj (2)	
model.audio_tower.output_proj.bias	[1 536]
model.audio_tower.output_proj.weight	[1 536, 1 024
model.language_model (6)	
model.language_model.embed_tokens.weight	[262 144, 2 5
model.language_model.embed_tokens_per_layer.weight	[262 144, 10
model.language_model.layers (42)	
model.language_model.layers.0 (10)	
model.language_model.layers.0.input_layernorm.weight	[2 560]
model.language_model.layers.0.layer_scalar	[1]
model.language_model.layers.0.mlp (3)	
model.language_model.layers.0.mlp.down_proj.weight	[2 560, 10 24
model.language_model.layers.0.mlp.gate_proj.weight	[10 240, 2 56
model.language_model.layers.0.mlp.up_proj.weight	[10 240, 2 56
model.language_model.layers.0.per_layer_input_gate.weight	[256, 2 560]
model.language_model.layers.0.per_layer_projection.weight	[2 560, 256]
model.language_model.layers.0.post_attention_layernorm.weight	[2 560]
model.language_model.layers.0.post_feedforward_layernorm.weight	[2 560]
model.language_model.layers.0.post_per_layer_input_norm.weight	[2 560]
model.language_model.layers.0.pre_feedforward_layernorm.weight	[2 560]
model.language_model.layers.0.self_attn (6)	
model.language_model.layers.0.self_attn.k_norm.weight	[256]
model.language_model.layers.0.self_attn.k_proj.weight	[512, 2 560]
model.language_model.layers.0.self_attn.o_proj.weight	[2 560, 2 048
model.language_model.layers.0.self_attn.q_norm.weight	[256]
model.language_model.layers.0.self_attn.q_proj.weight	[2 048, 2 560
model.language_model.layers.0.self_attn.v_proj.weight	[512, 2 560]
model.language_model.layers.1 (10)	

Tensors	Shape
model.language_model.layers.1.input_layernorm.weight	[2 560]
model.language_model.layers.1.layer_scalar	[1]
model.language_model.layers.1.mlp (3)	
model.language_model.layers.1.mlp.down_proj.weight	[2 560, 10 24
model.language_model.layers.1.mlp.gate_proj.weight	[10 240, 2 56
model.language_model.layers.1.mlp.up_proj.weight	[10 240, 2 56
model.language_model.layers.1.per_layer_input_gate.weight	[256, 2 560]
model.language_model.layers.1.per_layer_projection.weight	[2 560, 256]
model.language_model.layers.1.post_attention_layernorm.weight	[2 560]
model.language_model.layers.1.post_feedforward_layernorm.weight	[2 560]
model.language_model.layers.1.post_per_layer_input_norm.weight	[2 560]
model.language_model.layers.1.pre_feedforward_layernorm.weight	[2 560]
model.language_model.layers.1.self_attn (6)	
model.language_model.layers.1.self_attn.k_norm.weight	[256]
model.language_model.layers.1.self_attn.k_proj.weight	[512, 2 560]
model.language_model.layers.1.self_attn.o_proj.weight	[2 560, 2 048
model.language_model.layers.1.self_attn.q_norm.weight	[256]
model.language_model.layers.1.self_attn.q_proj.weight	[2 048, 2 560
model.language_model.layers.1.self_attn.v_proj.weight	[512, 2 560]
model.language_model.layers.2 (10)	
model.language_model.layers.3 (10)	
model.language_model.layers.4 (10)	
model.language_model.layers.5 (10)	
model.language_model.layers.6 (10)	
model.language_model.layers.7 (10)	
model.language_model.layers.8 (10)	
model.language_model.layers.9 (10)	
model.language_model.layers.10 (10)	
model.language_model.layers.11 (10)	
model.language_model.layers.12 (10)	
model.language_model.layers.13 (10)	
model.language_model.layers.14 (10)	
model.language_model.layers.15 (10)	
model.language_model.layers.16 (10)	
model.language_model.layers.17 (10)	
model.language_model.layers.18 (10)	
model.language_model.layers.19 (10)	
model.language_model.layers.20 (10)	

Tensors	Shape
model.language_model.layers.21 (10)	
model.language_model.layers.22 (10)	
model.language_model.layers.23 (10)	
model.language_model.layers.24 (10)	
model.language_model.layers.25 (10)	
model.language_model.layers.26 (10)	
model.language_model.layers.27 (10)	
model.language_model.layers.28 (10)	
model.language_model.layers.29 (10)	
model.language_model.layers.30 (10)	
model.language_model.layers.31 (10)	
model.language_model.layers.32 (10)	
model.language_model.layers.33 (10)	
model.language_model.layers.34 (10)	
model.language_model.layers.35 (10)	
model.language_model.layers.36 (10)	
model.language_model.layers.37 (10)	
model.language_model.layers.38 (10)	
model.language_model.layers.39 (10)	
model.language_model.layers.40 (10)	
model.language_model.layers.41 (10)	
model.language_model.norm.weight	[2 560]
model.language_model.per_layer_model_projection.weight	[10 752, 2 56
model.language_model.per_layer_projection_norm.weight	[256]
model.vision_tower (2)	
model.vision_tower.patch_embedder (2)	
model.vision_tower.patch_embedder.position_embedding_table	[2, 10 240,
model.vision_tower.patch_embedder.input_proj.weight	[768, 768]
model.vision_tower.encoder.layers (16)	
model.vision_tower.encoder.layers.0 (6)	
model.vision_tower.encoder.layers.0.input_layernorm.weight	[768]
model.vision_tower.encoder.layers.0.mlp (15)	
model.vision_tower.encoder.layers.0.mlp.down_proj.input_max	scalar
model.vision_tower.encoder.layers.0.mlp.down_proj.input_min	scalar
model.vision_tower.encoder.layers.0.mlp.down_proj.linear.weight	[768, 3 072]
model.vision_tower.encoder.layers.0.mlp.down_proj.output_max	scalar
model.vision_tower.encoder.layers.0.mlp.down_proj.output_min	scalar
model.vision_tower.encoder.layers.0.mlp.gate_proj.input_max	scalar

Tensors	Shape
model.vision_tower.encoder.layers.0.mlp.gate_proj.input_min	scalar
model.vision_tower.encoder.layers.0.mlp.gate_proj.linear.weight	[3 072, 768]
model.vision_tower.encoder.layers.0.mlp.gate_proj.output_max	scalar
model.vision_tower.encoder.layers.0.mlp.gate_proj.output_min	scalar
model.vision_tower.encoder.layers.0.mlp.up_proj.input_max	scalar
model.vision_tower.encoder.layers.0.mlp.up_proj.input_min	scalar
model.vision_tower.encoder.layers.0.mlp.up_proj.linear.weight	[3 072, 768]
model.vision_tower.encoder.layers.0.mlp.up_proj.output_max	scalar
model.vision_tower.encoder.layers.0.mlp.up_proj.output_min	scalar
model.vision_tower.encoder.layers.0.post_attention_layernorm.weight	[768]
model.vision_tower.encoder.layers.0.post_feedforward_layernorm.weight	[768]
model.vision_tower.encoder.layers.0.pre_feedforward_layernorm.weight	[768]
model.vision_tower.encoder.layers.0.self_attn (22)	
model.vision_tower.encoder.layers.0.self_attn.k_norm.weight	[64]
model.vision_tower.encoder.layers.0.self_attn.k_proj.input_max	scalar
model.vision_tower.encoder.layers.0.self_attn.k_proj.input_min	scalar
model.vision_tower.encoder.layers.0.self_attn.k_proj.linear.weight	[768, 768]
model.vision_tower.encoder.layers.0.self_attn.k_proj.output_max	scalar
model.vision_tower.encoder.layers.0.self_attn.k_proj.output_min	scalar
model.vision_tower.encoder.layers.0.self_attn.o_proj.input_max	scalar
model.vision_tower.encoder.layers.0.self_attn.o_proj.input_min	scalar
model.vision_tower.encoder.layers.0.self_attn.o_proj.linear.weight	[768, 768]
model.vision_tower.encoder.layers.0.self_attn.o_proj.output_max	scalar
model.vision_tower.encoder.layers.0.self_attn.o_proj.output_min	scalar
model.vision_tower.encoder.layers.0.self_attn.q_norm.weight	[64]
model.vision_tower.encoder.layers.0.self_attn.q_proj.input_max	scalar
model.vision_tower.encoder.layers.0.self_attn.q_proj.input_min	scalar
model.vision_tower.encoder.layers.0.self_attn.q_proj.linear.weight	[768, 768]
model.vision_tower.encoder.layers.0.self_attn.q_proj.output_max	scalar
model.vision_tower.encoder.layers.0.self_attn.q_proj.output_min	scalar
model.vision_tower.encoder.layers.0.self_attn.v_proj.input_max	scalar
model.vision_tower.encoder.layers.0.self_attn.v_proj.input_min	scalar
model.vision_tower.encoder.layers.0.self_attn.v_proj.linear.weight	[768, 768]
model.vision_tower.encoder.layers.0.self_attn.v_proj.output_max	scalar
model.vision_tower.encoder.layers.0.self_attn.v_proj.output_min	scalar
model.vision_tower.encoder.layers.1 (6)	
model.vision_tower.encoder.layers.2 (6)	
model.vision_tower.encoder.layers.3 (6)	

Tensors	Shape
model.vision_tower.encoder.layers.4 (6) >	
model.vision_tower.encoder.layers.5 (6) >	
model.vision_tower.encoder.layers.6 (6) >	
model.vision_tower.encoder.layers.7 (6) >	
model.vision_tower.encoder.layers.8 (6) >	
model.vision_tower.encoder.layers.9 (6) >	
model.vision_tower.encoder.layers.10 (6) >	
model.vision_tower.encoder.layers.11 (6) >	
model.vision_tower.encoder.layers.12 (6) >	
model.vision_tower.encoder.layers.13 (6) >	
model.vision_tower.encoder.layers.14 (6) >	
model.vision_tower.encoder.layers.15 (6) >	